

```
In [1]: import pandas as pd
```

```
In [2]: data=pd.read_csv("/home/placement/Downloads/fiat500.csv") #reads the file
```

```
In [3]: data.describe()
```

```
Out[3]:
```

	ID	engine_power	age_in_days	km	previous_owners	lat	lon	price
count	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000
mean	769.500000	51.904421	1650.980494	53396.011704	1.123537	43.541361	11.563428	8576.003901
std	444.126671	3.988023	1289.522278	40046.830723	0.416423	2.133518	2.328190	1939.958641
min	1.000000	51.000000	366.000000	1232.000000	1.000000	36.855839	7.245400	2500.000000
25%	385.250000	51.000000	670.000000	20006.250000	1.000000	41.802990	9.505090	7122.500000
50%	769.500000	51.000000	1035.000000	39031.000000	1.000000	44.394096	11.869260	9000.000000
75%	1153.750000	51.000000	2616.000000	79667.750000	1.000000	45.467960	12.769040	10000.000000
max	1538.000000	77.000000	4658.000000	235000.000000	4.000000	46.795612	18.365520	11100.000000

In [4]: data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1538 entries, 0 to 1537
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   ID                    1538 non-null   int64
1   model                 1538 non-null   object
2   engine_power          1538 non-null   int64
3   age_in_days           1538 non-null   int64
4   km                    1538 non-null   int64
5   previous_owners       1538 non-null   int64
6   lat                   1538 non-null   float64
7   lon                   1538 non-null   float64
8   price                 1538 non-null   int64
dtypes: float64(2), int64(6), object(1)
memory usage: 108.3+ KB
```

In [5]: data1=data.drop(['ID','lat','lon'],axis=1)

```
In [6]: data1
```

```
Out[6]:
```

	model	engine_power	age_in_days	km	previous_owners	price
0	lounge	51	882	25000	1	8900
1	pop	51	1186	32500	1	8800
2	sport	74	4658	142228	1	4200
3	lounge	51	2739	160000	1	6000
4	pop	73	3074	106880	1	5700
...
1533	sport	51	3712	115280	1	5200
1534	lounge	74	3835	112000	1	4600
1535	pop	51	2223	60457	1	7500
1536	lounge	51	2557	80750	1	5990
1537	pop	51	1766	54276	1	7900

1538 rows × 6 columns

```
In [7]: data2=pd.get_dummies(data1) #encodes the string into bits
```

```
In [8]: data2
```

```
Out[8]:
```

	engine_power	age_in_days	km	previous_owners	price	model_lounge	model_pop	model_sport
0	51	882	25000	1	8900	1	0	0
1	51	1186	32500	1	8800	0	1	0
2	74	4658	142228	1	4200	0	0	1
3	51	2739	160000	1	6000	1	0	0
4	73	3074	106880	1	5700	0	1	0
...
1533	51	3712	115280	1	5200	0	0	1
1534	74	3835	112000	1	4600	1	0	0
1535	51	2223	60457	1	7500	0	1	0
1536	51	2557	80750	1	5990	1	0	0
1537	51	1766	54276	1	7900	0	1	0

1538 rows × 8 columns

```
In [9]: data2.shape #shows num of rows and columns
```

```
Out[9]: (1538, 8)
```

```
In [10]: y=data2['price']#predicted value removed feom data frame  
x=data2.drop(['price'],axis=1)
```

```
In [11]: y #prices only will display
```

```
Out[11]: 0      8900
          1      8800
          2     4200
          3     6000
          4     5700
          ...
        1533    5200
        1534    4600
        1535    7500
        1536    5990
        1537    7900
        Name: price, Length: 1538, dtype: int64
```

```
In [12]: x #no prices
```

```
Out[12]:
```

	engine_power	age_in_days	km	previous_owners	model_lounge	model_pop	model_sport
0	51	882	25000	1	1	0	0
1	51	1186	32500	1	0	1	0
2	74	4658	142228	1	0	0	1
3	51	2739	160000	1	1	0	0
4	73	3074	106880	1	0	1	0
...
1533	51	3712	115280	1	0	0	1
1534	74	3835	112000	1	1	0	0
1535	51	2223	60457	1	0	1	0
1536	51	2557	80750	1	1	0	0
1537	51	1766	54276	1	0	1	0

1538 rows × 7 columns

```
In [13]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.33,random_state=42)#splitting into training a
```

```
In [14]: x_test.head(5)#shows the testing column values
```

```
Out[14]:
```

	engine_power	age_in_days	km	previous_owners	model_lounge	model_pop	model_sport
481	51	3197	120000	2	0	1	0
76	62	2101	103000	1	0	1	0
1502	51	670	32473	1	1	0	0
669	51	913	29000	1	1	0	0
1409	51	762	18800	1	1	0	0

```
In [15]: y_test.head(5)
```

```
Out[15]: 481    7900
76    7900
1502   9400
669    8500
1409   9700
Name: price, dtype: int64
```

```
In [16]: x_train.head(5)
```

```
Out[16]:
```

	engine_power	age_in_days	km	previous_owners	model_lounge	model_pop	model_sport
527	51	425	13111	1	1	0	0
129	51	1127	21400	1	1	0	0
602	51	2039	57039	1	0	1	0
331	51	1155	40700	1	1	0	0
323	51	425	16783	1	1	0	0

```
In [17]: y_train.head(5)
```

```
Out[17]: 527    9990
         129    9500
         602    7590
         331    8750
         323    9100
         Name: price, dtype: int64
```

```
In [18]: #linear regression starts
```

```
In [19]: from sklearn.linear_model import LinearRegression
         reg=LinearRegression() #creating object of linearregression
         reg.fit(x_train,y_train) #training and fitting LR object using training data
```

```
Out[19]: LinearRegression()
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [20]: ypred=reg.predict(x_test) #prediction of values(x_tesy*reg)
```

In [21]: ypred

```
Out[21]: array([ 5867.6503378 ,  7133.70142341,  9866.35776216,  9723.28874535,
 10039.59101162,  9654.07582608,  9673.14563045, 10118.70728123,
  9903.85952664,  9351.55828437, 10434.34963575,  7732.26255693,
  7698.67240131,  6565.95240435,  9662.90103518, 10373.20344286,
  9599.94844451,  7699.34400418,  4941.33017994, 10455.2719478 ,
 10370.51555682, 10391.60424404,  7529.06622456,  9952.37340054,
  7006.13845729,  9000.1780961 ,  4798.36770637,  6953.10376491,
  7810.39767825,  9623.80497535,  7333.52158317,  5229.18705519,
  5398.21541073,  5157.65652129,  8948.63632836,  5666.62365159,
  9822.1231461 ,  8258.46551788,  6279.2040404 ,  8457.38443276,
  9773.86444066,  6767.04074749,  9182.99904787, 10210.05195479,
  8694.90545226, 10328.43369248,  9069.05761443,  8866.7826029 ,
  7058.39787506,  9073.33877162,  9412.68162121, 10293.69451263,
 10072.49011135,  6748.5794244 ,  9785.95841801,  9354.09969973,
  9507.9444386 , 10443.01608254,  9795.31884316,  7197.84932877,
 10108.31707235,  7009.6597206 ,  9853.90699412,  7146.87414965,
  6417.69133992,  9996.97382441,  9781.18795953,  8515.83255277,
  8456.30006203,  6499.76668237,  7768.57829985,  6832.86406122,
  8347.96113362, 10439.02404036,  7356.43463051,  8562.56562053,
  8020.70555100, 10025.02571520,  7270.77100022,  8411.45004000]
```

In [22]: `from sklearn.metrics import r2_score#efficiency`
`r2_score(y_test,ypred) #ypred is predicted value`

Out[22]: 0.8415526986865394

In [23]: `from sklearn.metrics import mean_squared_error#to calculate rms value`

In [24]: `mean_squared_error(ypred,y_test)`

Out[24]: 581887.727391353


```
In [25]: results=pd.DataFrame(columns=['Price', 'Predicted'])
results['Price']=y_test
results['Predicted']=ypred
results=results.reset_index()
results['Id']=results.index
results.head(15)
```

```
Out[25]:
```

	index	Price	Predicted	Id
0	481	7900	5867.650338	0
1	76	7900	7133.701423	1
2	1502	9400	9866.357762	2
3	669	8500	9723.288745	3
4	1409	9700	10039.591012	4
5	1414	9900	9654.075826	5
6	1089	9900	9673.145630	6
7	1507	9950	10118.707281	7
8	970	10700	9903.859527	8
9	1198	8999	9351.558284	9
10	1088	9890	10434.349636	10
11	576	7990	7732.262557	11
12	965	7380	7698.672401	12
13	1488	6800	6565.952404	13
14	1432	8900	9662.901035	14

```
In [26]: results['final']=results.apply(lambda row:row.Price - row.Predicted,axis=1)
```

In [27]: results

Out[27]:


	index	Price	Predicted	Id	final
0	481	7900	5867.650338	0	2032.349662
1	76	7900	7133.701423	1	766.298577
2	1502	9400	9866.357762	2	-466.357762
3	669	8500	9723.288745	3	-1223.288745
4	1409	9700	10039.591012	4	-339.591012
...
503	291	10900	10032.665135	503	867.334865
504	596	5699	6281.536277	504	-582.536277
505	1489	9500	9986.327508	505	-486.327508
506	1436	6990	8381.517020	506	-1391.517020
507	575	10900	10371.142553	507	528.857447

508 rows × 5 columns

```
In [28]: import seaborn as hh
import matplotlib.pyplot as plt
```

```
In [29]: hh.lineplot(x='Id',y='Price',data=results.head(50))  
hh.lineplot(x='Id',y='Predicted',data=results.head(50))
```

```
Out[29]: <Axes: xlabel='Id', ylabel='Price'>
```



```
In [30]: #linear regression ends
```

```
In [31]: #ridge regression starts
```

```
In [32]: import warnings  
warnings.filterwarnings("ignore")
```

```
In [33]: from sklearn.model_selection import GridSearchCV  
from sklearn.linear_model import Ridge  
  
alpha = [1e-15, 1e-10, 1e-8, 1e-4, 1e-3, 1e-2, 1, 5, 10, 20, 30]  
  
ridge = Ridge()  
  
parameters = {'alpha': alpha}  
  
ridge_regressor = GridSearchCV(ridge, parameters)  
  
ridge_regressor.fit(x_train, y_train)
```

```
Out[33]: GridSearchCV(estimator=Ridge(),  
                      param_grid={'alpha': [1e-15, 1e-10, 1e-08, 0.0001, 0.001, 0.01, 1,  
                                             5, 10, 20, 30]}))
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [34]: ridge_regressor.best_params_
```

```
Out[34]: {'alpha': 30}
```

```
In [35]: ridge=Ridge(alpha=30)
```

```
In [36]: ridge.fit(x_train,y_train)  
y_pred_ridge=ridge.predict(x_test)
```

```
In [37]: from sklearn.metrics import mean_squared_error
Ridge_Error=mean_squared_error(y_pred_ridge,y_test)
Ridge_Error
```

Out[37]: 579521.7970897449

```
In [38]: from sklearn.metrics import r2_score
r2_score(y_test,y_pred_ridge)
```

Out[38]: 0.8421969385523054

```
In [39]: results=pd.DataFrame(columns=['actual','Predicted'])
results['actual']=y_test
results['Predicted']=y_pred_ridge
results=results.reset_index()
results['Id']=results.index
results.head(10)
```


Out[39]:

	index	actual	Predicted	Id
0	481	7900	5869.741155	0
1	76	7900	7149.563327	1
2	1502	9400	9862.785355	2
3	669	8500	9719.283532	3
4	1409	9700	10035.895686	4
5	1414	9900	9650.311090	5
6	1089	9900	9669.183317	6
7	1507	9950	10115.128380	7
8	970	10700	9900.241944	8
9	1198	8999	9347.080772	9

```
In [40]: import seaborn as hh
import matplotlib.pyplot as plt
```

```
In [41]: hh.lineplot(x='Id',y='actual',data=results.head(50))  
hh.lineplot(x='Id',y='Predicted',data=results.head(50))
```

Out[41]: <Axes: xlabel='Id', ylabel='actual'>



```
In [42]: #ridge regression ends
```

```
In [43]: #elastic model starts
```

```
In [44]: from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import ElasticNet
elastic = ElasticNet()
parameters = {'alpha': [1e-15, 1e-10, 1e-8, 1e-4, 1e-3, 1e-2, 1, 5, 10, 20]}
elastic_regressor = GridSearchCV(elastic, parameters)
elastic_regressor.fit(x_train, y_train)
```

```
Out[44]: GridSearchCV(estimator=ElasticNet(),
                      param_grid={'alpha': [1e-15, 1e-10, 1e-08, 0.0001, 0.001, 0.01, 1,
                                             5, 10, 20]})
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [45]: elastic_regressor.best_params_
```

```
Out[45]: {'alpha': 0.01}
```

```
In [46]: elastic=ElasticNet(alpha=30)
```

```
In [47]: elastic.fit(x_train,y_train)
y_pred_elastic=elastic.predict(x_test)
```

```
In [48]: from sklearn.metrics import mean_squared_error
ElasticNet_Error=mean_squared_error(y_pred_elastic,y_test)
ElasticNet_Error
```

```
Out[48]: 580334.1755711779
```

```
In [49]: from sklearn.metrics import r2_score
r2_score(y_test,y_pred_elastic)
```

```
Out[49]: 0.8419757289065801
```

```
In [50]: results=pd.DataFrame(columns=['actual','Predicted'])
results['actual']=y_test
results['Predicted']=y_pred_elastic
results=results.reset_index()
results['Id']=results.index
results.head(10)
```

```
Out[50]:
```


	index	actual	Predicted	Id
0	481	7900	5999.772939	0
1	76	7900	7307.696255	1
2	1502	9400	9811.206661	2
3	669	8500	9664.419998	3
4	1409	9700	9983.473801	4
5	1414	9900	9597.210309	5
6	1089	9900	9614.618393	6
7	1507	9950	10063.607164	7
8	970	10700	9848.342378	8
9	1198	8999	9288.542203	9

```
In [51]: import seaborn as hh
import matplotlib.pyplot as plt
```



```
In [52]: hh.lineplot(x='Id',y='actual',data=results.head(50))  
hh.lineplot(x='Id',y='Predicted',data=results.head(50))
```

```
Out[52]: <Axes: xlabel='Id', ylabel='actual'>
```



```
In [ ]: #elastic model ends
```